

Using Service Mesh Technology within a Microservice architecture design. Environment Setup

by

Eamon Foy

This thesis has been submitted in partial fulfillment for the
degree of Master of Science in MSc in Cloud Computing

in the

Faculty of Engineering and Science
Department of Computer Science

April 2020

Contents

Environment Setup	ii
Building the base Kubernetes system	ii
Curl Install	iii
Install Virtual box	iii
Install vagrant	iii
Install Ansible	iv
Install a kubernetes cluster, with 1 master and 3 nodes	iv
Installing kubectl on host	vii
Copy the Kubernetes config	viii
Check is kubectl working	viii
Deploy a test app to check if the cluster is working	x
Install the kubernetes Dashboard	xii
Install a Load balancer (metallb)	xv
Background	xv
To install MetalLB, apply the following manifest:	xvi
Testing the load balancer	xviii
Install NFS Server on the host	xx
Install Dynamic Storage Volume Creation	xxi
Install HA Proxy	xxiii
HA Proxy Configuration	xxiv
Allow other machines to access the cluster:	xxvi
Finding Your Network Details	xxvi
Troubleshooting Issues:	xxviii
Useful commands	xxix
Vagrant commands	xxix
Taking vagrant snapshots	xxix
Suspending the vagrant virtual machines	xxx
Install Helm	xxx
Install Istio []	xxxi
Install Kubernetes Cert Manager	xxxiv
Enable Remote access to the telemetry Addons	xxxv
Delete istio	xxxv

Environment Setup

The purpose of this document is to provide a development guide which outlines how to install and configure all the different components which went into making this project. Starting with installing Virtualbox, and ansible, provisioning the cluster with vagrant. Install and configure an NFS storage to act as permanent storage for databases and Elk stack. Providing bare-metal load balancing with MetalLb. Installing all the necessary components for Istio. Installing and configuring HAProxy to expose all the necessary endpoint.

Building the base Kubernetes system

The base system consists of the following installation steps:

1. [Install curl](#)
2. [Install Virtual box](#)
3. [Install Vagrant](#)
4. [Install Ansible](#)
5. [Install a kubernetes cluster, with 1 master and 3 nodes](#)
6. [Install the kubernetes Dashboard](#)
7. [Install a Load balancer \(metallb\)](#)
8. [Install Helm](#)
9. [Install Dynamic Storage Volume Creation in Kubernetes](#)
10. [Install NFS Server on the host](#)
11. [Install Istio](#)
12. [Allow other machines to access the cluster](#)
13. [Useful commands](#)

Curl Install

Installing curl through the command line

```
1 sudo apt install curl
```

Install Virtual box

Installing VirtualBox Through the Command line

```
1 sudo add-apt-repository multiverse && sudo apt-get update
2 sudo apt install virtualbox
```

Install the Extension Package

```
1 sudo apt install virtualbox-ext-pack
```

Check the virtual box version, it should be 6.0 or above

```
1 vboxmanage --version
2 $ 6.0.6_Ubuntu129722
```

Run virtual box with

```
1 virtualbox
```

Install vagrant

Start by updating the package list with:

```
1 sudo apt update
```

Download the Vagrant package using the following curl command:

```
1 curl -O
    https://releases.hashicorp.com/vagrant/2.2.6/vagrant_2.2.6_x86_64.deb
```

Once the .deb file is downloaded, install it by typing:

```
1 sudo apt install ./vagrant_2.2.6_x86_64.deb
```

Verify Vagrant installation To verify that the installation was successful, run the following command which prints the Vagrant version:

```
1 vagrant --version
```

The output should look something like this:

```
1 Vagrant 2.2.6
```

Install Ansible

Installing Ansible on the Command line

```
1 sudo apt update
2 sudo apt install software-properties-common
3 sudo apt-add-repository --yes --update ppa:ansible/ansible
4 sudo apt install ansible -y
```

Check the version of Ansible

```
1 $ ansible --version
2
3 ansible 2.9.6
4   config file = /etc/ansible/ansible.cfg
5   configured module search path =
6     [u'/home/eamonfoy/.ansible/plugins/modules',
7      u'/usr/share/ansible/plugins/modules']
6   ansible python module location =
7     /usr/lib/python2.7/dist-packages/ansible
7   executable location = /usr/bin/ansible
8   python version = 2.7.17 (default, Nov  7 2019, 10:07:09) [GCC 7.4.0]
```

Install a kubernetes cluster, with 1 master and 3 nodes

This section assumes you have successfully downloaded source code. From the root of the source code folder follow the following commands

```
1 cd 1_base_system_build/kubernetes-vagrant_cluster
```

This section assumes you have successfully downloaded source code. From the root of the source code folder follow the following commands

```
1 vagrant up
```

This command will take a while depending on your internet speed and also the capacity of your machine. For me it was around 10 minutes.

This section is inspired by a tutorial found [here](#)

We are going to use vagrant to set up the following vm's which will contain the kubernetes cluster:

vm(s)	Description	IP	RAM	vCPU
K8S-M-1	Kubernetes Master	192.168.50.11	4Gb	4
K8S-N-1	Kubernetes Node 1	192.168.50.12	16GB	4
K8S-N-2	Kubernetes Node 2	192.168.50.13	16GB	4
K8S-N-3	Kubernetes Node 3	192.168.50.14	16GB	4

Virtualbox NAT network setting are:

Description	IP
host	10.0.0.45
vboxnet0 virtual iface	192.168.50.1

Virtual box creates routes which you will need to make not of

```
1 $ route
2 Kernel IP routing table
3 Destination      Gateway            Genmask           Flags Metric Ref    Use
   Iface
4 default          _gateway          0.0.0.0           UG    100    0      0
   eno1
5 10.0.0.0          0.0.0.0           255.255.255.0    U     100    0      0
   eno1
```

6	link-local	0.0.0.0	255.255.0.0	U	1000	0	0
	eno1						
7	192.168.50.0	0.0.0.0	255.255.255.0	U	0	0	0
	vboxnet0						

To ssh in to the master node use the following: “bash ssh vagrant@192.168.50.11

Note when it asks for a password type : vagrant “bash

Now I am on the kubernetes master

```

1 ssh vagrant@192.168.50.11
2 The authenticity of host '192.168.50.11 (192.168.50.11)' can't be
   established.
3 ECDSA key fingerprint is
   SHA256:7ok3hLsAT6RMxwYkoISSQ5XltpgIng6eQ6ZJk/bkOMY.
4 Are you sure you want to continue connecting (yes/no)? yes
5 Warning: Permanently added '192.168.50.11' (ECDSA) to the list of known
   hosts.
6 vagrant@192.168.50.11's password:
7 Welcome to Ubuntu 18.04.4 LTS (GNU/Linux 4.15.0-76-generic x86_64)
8
9 * Documentation:  https://help.ubuntu.com
10 * Management:    https://landscape.canonical.com
11 * Support:       https://ubuntu.com/advantage
12
13 System information as of Sat Feb  8 23:46:48 UTC 2020
14
15 System load:  1.23           Users logged in:      0
16 Usage of /:   5.7% of 61.80GB IP address for eth0:   10.0.2.15
17 Memory usage: 40%           IP address for eth1:   192.168.50.11
18 Swap usage:   0%            IP address for docker0: 172.17.0.1
19 Processes:    146           IP address for tunl0:   192.168.116.0
20
21 * Multipass 1.0 is out! Get Ubuntu VMs on demand on your Linux, Windows
   or

```

```
22 Mac. Supports cloud-init for fast, local, cloud devops simulation.
23
24     https://multipass.run/
25
26 9 packages can be updated.
27 9 updates are security updates.
28
29
30
31 This system is built by the Bento project by Chef Software
32 More information can be found at https://github.com/chef/bento
33 Last login: Sat Feb  8 21:56:40 2020 from 10.0.2.2
```

While we are on the kubernetes master we can check the status of the cluster:

```
1 kubectl get all
```

Installing kubectl on host

This will enable us to administer the Kubernetes Cluster from host and not needing to log into the master.

Run the following commands from the host machine:

```
1
2 sudo apt-get update && sudo apt-get install -y apt-transport-https
3
4 curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | sudo
    apt-key add -
5
6 echo "deb https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
    -a /etc/apt/sources.list.d/kubernetes.list
7
8 sudo apt-get update
9
10 sudo apt-get install -y kubectl
```


Copy the Kubernetes config

This series of steps will copy the Kubernetes config to your local home .kube dir

Create the configuration directory

```
1 $ mkdir -p ~/.kube
```

Find the SSH port of the k8s-m-1 server

```
1 $ vagrant port k8s-m-1
2
3 The forwarded ports for the machine are listed below. Please note that
4 these values may differ from values configured in the Vagrantfile if the
5 provider supports automatic port collision detection and resolution.
6     22 (guest) => 2222 (host)
```

Copy the file using scp (ssh password is vagrant)

```
1 $ scp -P 2222 vagrant@127.0.0.1:/home/vagrant/.kube/config
   ~/.kube/config
2
3 vagrant@127.0.0.1's password: vagrant
4 config
   100% 5449   118.7KB/s   00:00
```

Check is kubectl working

List the Kubernetes cluster nodes using kubectl from your development host:

```
1 $ kubectl cluster-info
2
3 Kubernetes master is running at https://192.168.50.11:6443
4 KubeDNS is running at
   https://192.168.50.11:6443/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy
```

Check the status of the nodes

```
1 $ kubectl get nodes --all-namespaces
```

2	NAME	STATUS	ROLES	AGE	VERSION
3	k8s-m-1	Ready	master	49m	v1.17.2
4	k8s-n-1	Ready	<none>	46m	v1.17.2
5	k8s-n-2	Ready	<none>	43m	v1.17.2

Check the status of all the pods

1	\$ kubectl get pods --all-namespaces				
2	NAMESPACE	NAME	READY	STATUS	
		RESTARTS	AGE		
3	kube-system	calico-kube-controllers-6895d4984b-jbg6r	1/1		
		Running	0	52m	
4	kube-system	calico-node-pftj4	1/1		
		Running	0	52m	
5	kube-system	calico-node-vdshc	1/1		
		Running	0	6m	
6	kube-system	calico-node-w7s7v	1/1		
		Running	0	49m	
7	kube-system	coredns-6955765f44-6qk9z	1/1		
		Running	0	52m	
8	kube-system	coredns-6955765f44-tql14	1/1		
		Running	0	52m	
9	kube-system	etcd-k8s-m-1	1/1		
		Running	0	52m	
10	kube-system	kube-apiserver-k8s-m-1	1/1		
		Running	0	52m	
11	kube-system	kube-controller-manager-k8s-m-1	1/1		
		Running	0	52m	
12	kube-system	kube-proxy-6f692	1/1		
		Running	0	52m	
13	kube-system	kube-proxy-6fhzt	1/1		
		Running	0	46m	
14	kube-system	kube-proxy-h72z6	1/1		
		Running	0	49m	

15	kube-system	kube-scheduler-k8s-m-1	1/1
	Running	0	52m

Deploy a test app to check if the cluster is working

This section assumes you have successfully downloaded source code.

From the root of the source code folder follow the following commands

```
1 cd 1_base_system_build/kubernetes-vagrant_cluster
```

Deploy nginx to the cluster

```
1 $ kubectl apply -f nginx.yaml
2 deployment.apps/nginx-deployment created
3 service/nginx-service-np created
```

We need to check the application has been deployed correctly.

To do this we will now check the two nodes using curl to check that nginx got deployed correctly.

Check Node 1

```
1 $ curl http://192.168.50.12:30000/
2
3 Hostname: nginx-deployment-569b77699b-lvstv
4
5 Pod Information:
6   -no pod information available-
7
8 Server values:
9   server_version=nginx: 1.13.3 - lua: 10008
10
11 Request Information:
12   client_address=192.168.122.0
13   method=GET
14   real path=/
```

```
15     query=
16     request_version=1.1
17     request_uri=http://192.168.50.12:8080/
18
19 Request Headers:
20     accept=/*/*
21     host=192.168.50.12:30000
22     user-agent=curl/7.64.0
23
24 Request Body:
25     -no body in request-
```

Check Node 2

```
1 $ curl http://192.168.50.13:30000/
2
3 Hostname: nginx-deployment-569b77699b-lvstv
4
5 Pod Information:
6     -no pod information available-
7
8 Server values:
9     server_version=nginx: 1.13.3 - lua: 10008
10
11 Request Information:
12     client_address=192.168.122.0
13     method=GET
14     real path=/
15     query=
16     request_version=1.1
17     request_uri=http://192.168.50.13:8080/
18
19 Request Headers:
20     accept=/*/*
```

```

21     host=192.168.50.13:30000
22     user-agent=curl/7.64.0
23
24 Request Body:
25     -no body in request-
```

Install the kubernetes Dashboard

The Kubernetes Dashboard provides a web-based user interface to deploy applications, troubleshoot and manage resources. The same functionality is provided through the command line tools but under a very nice web application with charts and beautiful screens.

To deploy the Web UI (Dashboard) or Kubernetes Dashboard run the following command:

```

1 kubectl apply -f
    https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0-beta8/aio/deploy/reco
```

The deployment file will publish the Kubernetes Dashboard using a ClusterIP service as shown below using TargetPort 8443:

```

1 $ kubectl -n kubernetes-dashboard describe service kubernetes-dashboard
```

response:

```

1 Name:                kubernetes-dashboard
2 Namespace:           kubernetes-dashboard
3 Labels:              k8s-app=kubernetes-dashboard
4 Annotations:         kubectl.kubernetes.io/last-applied-configuration:
5
6     {"apiVersion":"v1","kind":"Service","metadata":{"annotations":{},"labels":{"k8s-ap
7 Selector:            k8s-app=kubernetes-dashboard
8 Type:                ClusterIP
9 IP:                  10.109.172.25
10 Port:                <unset> 443/TCP
11 TargetPort:          8443/TCP
```

11	Endpoints:	192.168.122.129:8443
12	Session Affinity:	None
13	Events:	<none>

In order to access the Kubernetes Dashboard from our workstation, a NodePort will be created to publish the kubernetes-dashboard following the Publish an Application Outside Kubernetes Cluster instructions.

The file kubernetes-dashboard-service-np.yaml:

Creates an admin-user Assigns the cluster-admin role Creates a new NodePort service that publishes TargetPort 8443 as NodePort 30002:

```
1 ---
2 apiVersion: v1
3 kind: ServiceAccount
4 metadata:
5   name: admin-user
6   namespace: kubernetes-dashboard
7 ---
8 apiVersion: rbac.authorization.k8s.io/v1
9 kind: ClusterRoleBinding
10 metadata:
11   name: admin-user
12 roleRef:
13   apiGroup: rbac.authorization.k8s.io
14   kind: ClusterRole
15   name: cluster-admin
16 subjects:
17 - kind: ServiceAccount
18   name: admin-user
19   namespace: kubernetes-dashboard
20 ---
21 kind: Service
22 apiVersion: v1
23 metadata:
```

```
24 namespace: kubernetes-dashboard
25 name: kubernetes-dashboard-service-np
26 labels:
27   k8s-app: kubernetes-dashboard
28 spec:
29   type: NodePort
30   ports:
31   - port: 8443
32     nodePort: 30002
33     targetPort: 8443
34     protocol: TCP
35   selector:
36     k8s-app: kubernetes-dashboard
```

Apply the changes:

```
1 kubectl apply -f kubernetes-dashboard-service-np.yaml
```

Obtain an authentication token to use on the Kubernetes Dashboard authentication realm:

```
1 kubectl -n kubernetes-dashboard describe secret $(kubectl -n
  kubernetes-dashboard get secret | grep admin-user | awk '{print $1}')
```

The command will print out the token, yours will be different than mine

```
1 kubectl -n kubernetes-dashboard describe secret $(kubectl -n
  kubernetes-dashboard get secret | grep admin-user | awk '{print $1}')
```

```
2 Name:          admin-user-token-vjh27
3 Namespace:     kubernetes-dashboard
4 Labels:        <none>
5 Annotations:   kubernetes.io/service-account.name: admin-user
6               kubernetes.io/service-account.uid:
               ab754545-5c40-463a-9ce7-0f8d19279d23
7
8 Type:          kubernetes.io/service-account-token
```

```
9
10 Data
11 ====
12 ca.crt:      1025 bytes
13 namespace:   20 bytes
14 token:
    eyJhbGciOiJSUzI1NiIsImtpZCI6Ikk4xeGprVGNuMzBWRGtLamhVRGRYZl9YaTY3TWNBNTlByQWFaei1YeU
```

Copy the token part on to the clipboard for later

Access the Kubernetes Dashboard using the URL <https://192.168.50.11:30002/#/login>

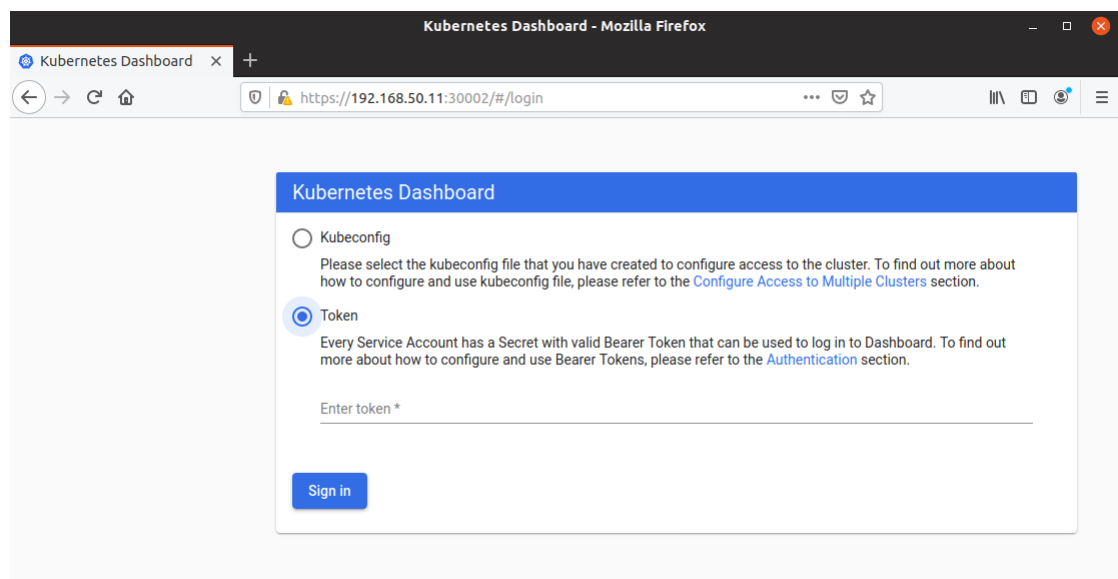


FIGURE 1: login

Using the token copied to the clipboard, paste in and click sign in. This will allow you to login in to the Dashboard:

Install a Load balancer (metallb)

Background

Because Kubernetes does not offer an implementation of network load-balancers (Services of type LoadBalancer) for bare metal clusters.

The implementations of Network LB that Kubernetes does ship with are all glue code that calls out to various IaaS platforms (GCP, AWS, Azure...). If you're not running

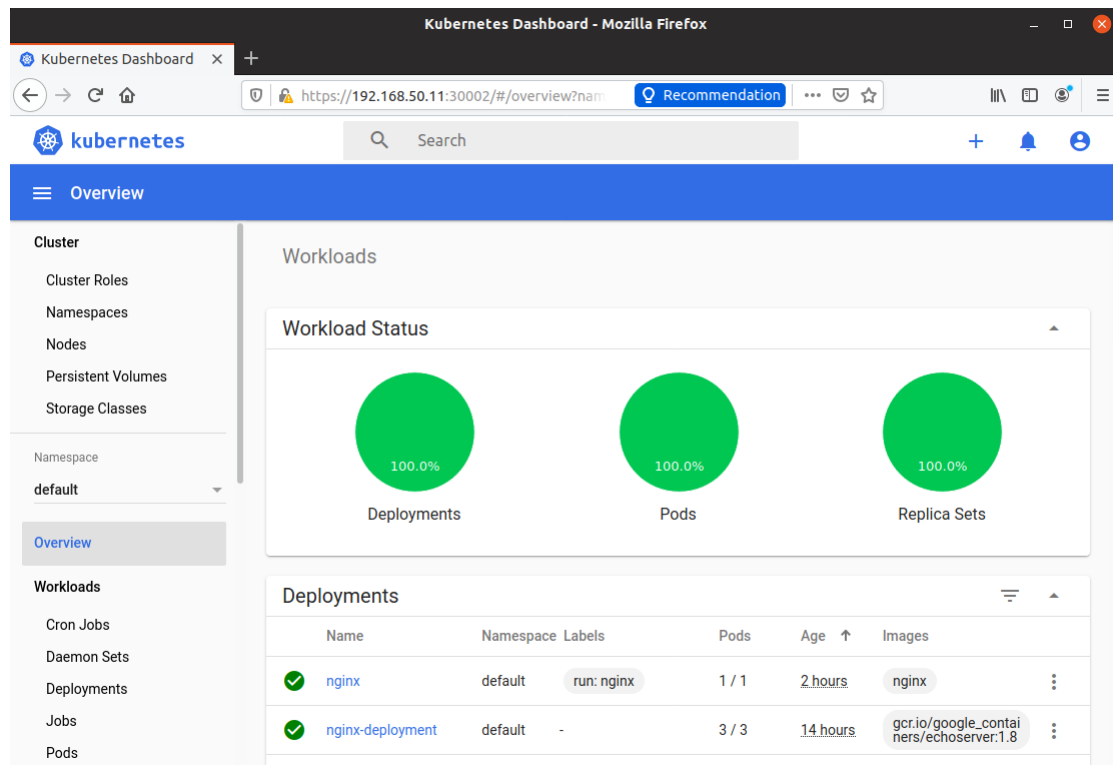


FIGURE 2: login

on a supported IaaS platform (GCP, AWS, Azure...), LoadBalancers will remain in the “pending” state indefinitely when created.

Bare metal cluster operators are left with two lesser tools to bring user traffic into their clusters, “NodePort” and “externalIPs” services. Both of these options have significant downsides for production use, which makes bare metal clusters second class citizens in the Kubernetes ecosystem.

MetalLB aims to redress this imbalance by offering a Network LB implementation that integrates with standard network equipment, so that external services on bare metal clusters also “just work” as much as possible.

To install **MetalLB**, apply the following manifest:

```
1 kubectl apply -f
    https://raw.githubusercontent.com/google/metallb/v0.8.3/manifests/metallb.yaml
```

To configure metal lb we need to find out more about the addressing which is being used within the cluster:

```
1 $ kubectl get nodes -o wide
```

2	NAME	STATUS	ROLES	AGE	VERSION	INTERNAL-IP	EXTERNAL-IP
	OS-IMAGE		KERNEL-VERSION		CONTAINER-RUNTIME		
3	k8s-m-1	Ready	master	12h	v1.17.2	192.168.50.11	<none>
	Ubuntu 18.04.4 LTS		4.15.0-76-generic		docker://19.3.5		
4	k8s-n-1	Ready	<none>	12h	v1.17.2	192.168.50.12	<none>
	Ubuntu 18.04.4 LTS		4.15.0-76-generic		docker://19.3.5		
5	k8s-n-2	Ready	<none>	12h	v1.17.2	192.168.50.13	<none>
	Ubuntu 18.04.4 LTS		4.15.0-76-generic		docker://19.3.5		

Metallb is configured using the following command

Note: - The ip **addresss** range below will need to be changed depending on you kubernetes deployment. The previous command will advise you on which network range to use

```

1 cat <<EOF | kubectl create -f -
2 apiVersion: v1
3 kind: ConfigMap
4 metadata:
5   namespace: metallb-system
6   name: config
7 data:
8   config: |
9     address-pools:
10    - name: default
11      protocol: layer2
12      addresses:
13    - 192.168.50.240-192.168.50.250
14 EOF

```

**** ONLY REQUIRED IF YOUR RUN IN TO BOTHER *** Check that the configuration is set by using the following command

```
1 kubectl get configmap -n metallb-system config -o json
```

Just in case you need to delete this configuration and strat over, here is the command:

```
1 kubectl delete configmap -n metallb-system config
```

Testing the load balancer

Deploy a simple nginx container

```
1 kubectl run nginx --image nginx
```

Create a loadbalancer service to expose the service on port 80

```
1 kubectl expose deploy nginx --port 80 --type LoadBalancer
```

list all to see

```
1 kubectl get all -o wide
```

See below the service is exposed on ip 192.168.50.240 port 80:32049

```
1 NAME                                READY   STATUS    RESTARTS
   AGE   IP              NODE     NOMINATED NODE   READINESS GATES
2 pod/nginx-6db489d4b7-6mbm2          1/1     Running   0
   81s   192.168.122.133 k8s-n-2   <none>           <none>
3 pod/nginx-deployment-569b77699b-44hs4 1/1     Running   1
   12h   192.168.122.2   k8s-n-1   <none>           <none>
4 pod/nginx-deployment-569b77699b-6t4z7 1/1     Running   1
   12h   192.168.122.132 k8s-n-2   <none>           <none>
5 pod/nginx-deployment-569b77699b-lvstv 1/1     Running   1
   12h   192.168.122.131 k8s-n-2   <none>           <none>
6
7 NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP
   PORT(S)          AGE    SELECTOR
8 service/kubernetes  ClusterIP           10.96.0.1       <none>
   443/TCP          13h    <none>
9 service/nginx      LoadBalancer       10.104.202.13
   192.168.50.240   80:32049/TCP       27s    run=nginx
10 service/nginx-service-np  NodePort           10.103.144.87   <none>
   8082:30000/TCP   12h    app=nginx
11
12 NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
   CONTAINERS   IMAGES                               SELECTOR
```

```

13 deployment.apps/nginx          1/1      1          1          81s
    nginx          nginx          run=nginx
14 deployment.apps/nginx-deployment 3/3      3          3          12h
    my-echo        gcr.io/google_containers/echoserver:1.8  app=nginx
15
16 NAME                                DESIRED    CURRENT    READY
    AGE      CONTAINERS    IMAGES
    SELECTOR
17 replicaset.apps/nginx-6db489d4b7    1          1          1
    81s      nginx        nginx
    pod-template-hash=6db489d4b7,run=nginx
18 replicaset.apps/nginx-deployment-569b77699b 3          3          3
    12h      my-echo        gcr.io/google_containers/echoserver:1.8
    app=nginx,pod-template-hash=569b77699b

```

Test that everything is working by using curl as follows:

```
1 $ curl 192.168.50.240:80
```

Response will be as follows:

```

1 <!DOCTYPE html>
2 <html>
3 <head>
4 <title>Welcome to nginx!</title>
5 <style>
6     body {
7         width: 35em;
8         margin: 0 auto;
9         font-family: Tahoma, Verdana, Arial, sans-serif;
10    }
11 </style>
12 </head>
13 <body>
14 <h1>Welcome to nginx!</h1>

```

```
15 <p>If you see this page, the nginx web server is successfully installed
    and
16 working. Further configuration is required.</p>
17
18 <p>For online documentation and support please refer to
19 <a href="http://nginx.org/">nginx.org</a>.<br/>
20 Commercial support is available at
21 <a href="http://nginx.com/">nginx.com</a>.</p>
22
23 <p><em>Thank you for using nginx.</em></p>
24 </body>
25 </html>
```

Once your done with nginx you can remove as follows

```
1 kubectl delete deployment,pod,service nginx
```

Install NFS Server on the host

Install the Ubuntu needed packages

```
1 sudo apt install -y nfs-kernel-server
```

Make the neccessary directories and ownerships

```
1 sudo mkdir -p /mnt/vagrant-kubernetes
2 sudo mkdir -p /mnt/vagrant-kubernetes/data
3 sudo mkdir -p /mnt/vagrant-kubernetes/aodb
4 sudo mkdir -p /mnt/vagrant-kubernetes/flight
5 sudo mkdir -p /mnt/vagrant-kubernetes/weather
6 sudo mkdir -p /mnt/vagrant-kubernetes/elastic
7 sudo chown nobody:nogroup /mnt/vagrant-kubernetes
8 sudo chmod 777 /mnt/vagrant-kubernetes
```

Edit the `/etc/exports` file to add the exported local directory and limit the share to the CIDR `192.168.50.0/24` of our Kubernetes Vagrant Cluster VirtualBox machines.

```
1 /mnt/vagrant-kubernetes
   *(rw, sync, no_subtree_check, insecure, no_root_squash)
```

Start the nfs server

```
1 sudo exportfs -a
2 sudo systemctl restart nfs-kernel-server
3 sudo exportfs -v
```

This is the response you should see

```
1 /mnt/vagrant-kubernetes
2 *(rw, wdelay, insecure, no_root_squash, no_subtree_check, sec=sys, rw, insecure, no_root_squa
```

showmount -e

Install Dynamic Storage Volume Creation

The youtube video [Dynamically provision NFS persistent volumes in Kubernetes](#) is my inspiration

First of all make sure your in the following directory

```
1 cd
   01_system_build/kubernetes-vagrant_cluster/dynamic_nfs_persistent_volume_creation
```

This directory holds all the yamls necessary to install dynamic NFS storage provisioning

- class.yaml
- default-sc.yaml
- deployment.yaml
- rbac.yaml

We are going to create create the following items - ServiceAccount - Role - RoleBinding - ClusterRole - ClusterRoleBinding

by issueing the following command

```
1 kubectl create -f rbac.yaml
```

Next we are going to create the StorageClass and we are going to call this **managed-nfs-storage**. This is created by the following command:

```
1 kubectl create -f class.yaml
```

Next we will create a deployment for the nfs client provisioner.

Make sure to update the “server” and “path” within the deployment.yml to point to your own nfs server setup and additionally NFS_SERVER and NFS_PATH.

The following **deployment.yaml** represents my own setup and will be different in every situation.

```
1 kind: Deployment
2 apiVersion: apps/v1
3 metadata:
4   name: nfs-client-provisioner
5 spec:
6   replicas: 1
7   strategy:
8     type: Recreate
9   selector:
10    matchLabels:
11      app: nfs-client-provisioner
12   template:
13     metadata:
14       labels:
15         app: nfs-client-provisioner
16     spec:
17       serviceAccountName: nfs-client-provisioner
18       containers:
19         - name: nfs-client-provisioner
20           image: quay.io/external_storage/nfs-client-provisioner:latest
21           volumeMounts:
22             - name: nfs-client-root
23               mountPath: /persistentvolumes
24           env:
```

```
25         - name: PROVISIONER_NAME
26           value: example.com/nfs
27         - name: NFS_SERVER
28           value: 10.0.0.23
29         - name: NFS_PATH
30           value: /mnt/vagrant-kubernetes/data
31     volumes:
32       - name: nfs-client-root
33         nfs:
34           server: 10.0.0.23
35           path: /mnt/vagrant-kubernetes/data
```

By issuing the following command we will create the deployment:

```
1 kubectl create -f deployment.yaml
```

Install HA Proxy

This is the install process for HAProxy 1.7

```
1 sudo add-apt-repository ppa:vbernat/haproxy-1.7
```

Confirm adding the new PPA by pressing the Enter key.

Next, update your sources list.

```
1 sudo apt update
```

Then install HAProxy as you normally would.

```
1 sudo apt install -y haproxy
```

Afterwards, you can double check the installed version number with the following command.

```
1
2 haproxy -v
```



```
1 HA-Proxy version 1.7.8-1ppa1~xenial 2017/07/09
2 Copyright 2000-2017 Willy Tarreau <willy@haproxy.org>
```

The installation is then complete. Continue below with the instructions for how to configuring the load balancer to redirect requests to your web servers.

HA Proxy Configuration

Once installed HAProxy should already have a template for configuring the load balancer. Open the configuration file, for example, using nano with the command underneath.

```
1 sudo nano /etc/haproxy/haproxy.cfg
```

Add the following sections to the end of the file. Replace the with whatever you want to call your servers on the statistics page and the with the private IPs for the servers you wish to direct the web traffic to. You can check the private IPs at your UpCloud Control Panel and Private network tab under Network menu.

```
1 frontend http_front
2     bind *:80
3     stats uri /haproxy?stats
4     default_backend http_back
5
6 backend http_back
7     balance roundrobin
8     server <server1 name> <private IP 1>:80 check
9     server <server2 name> <private IP 2>:80 check
```

Here is my specific configuration

```
1 defaults
2     timeout connect 5000
3     timeout client 50000
4     timeout server 50000
5
6 #####
```

```
7 # Proxy for accessing the web front end and rest endpoints
8 #####
9 frontend http_front
10     bind *:81
11     stats uri /haproxy?stats
12     default_backend http_back
13
14 backend http_back
15     balance roundrobin
16     server node1 192.168.50.240:8080 check
17     server node1 192.168.50.240:8080 check
18
19 #####
20 # Proxy for accessing the kubernetes dashboard
21 #####
22 frontend http_front2
23     bind *:30002
24     stats uri /haproxy?stats
25     default_backend http_back2
26
27 backend http_back2
28     balance roundrobin
29     server node1 192.168.50.11:30002 check
30     server node1 192.168.50.11:30002 check
31
32 #####
33 # Proxy for accessing the Kibana dashboard
34 #####
35 frontend http_front3
36     bind *:5601
37     stats uri /haproxy?stats
38     default_backend http_back3
39
```

```
40 backend http_back3
41     balance roundrobin
42     server node1 192.168.50.241:5601 check
43     server node1 192.168.50.241:5601 check
```

After making the configurations, save the file and restart HAProxy with the next command.

```
1 sudo systemctl restart haproxy
```

If you get any errors or warnings at startup, check the configuration for any mistypes and then try restarting again.

Allow other machines to access the cluster:

First of all find out which machines you would like to provides access to and get their IP address in my case I have one machine which I would like to give access

Finding Your Network Details

To get the details of your own systems, begin by finding your network interfaces. You can find the interfaces on your machines and the addresses associated with them by typing:

```
1 ip -4 addr show scope global
```

Output:

```
1
2
3 3: enp1s0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel
   state UP group default qlen 1000
4   inet 10.0.0.23/24 brd 10.0.0.255 scope global dynamic noprefixroute
   enp1s0
5       valid_lft 47753sec preferred_lft 47753sec
6 4: docker0: <NO-CARRIER,BROADCAST,MULTICAST,UP> mtu 1500 qdisc noqueue
   state DOWN group default
7   inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
```

```

8     valid_lft forever preferred_lft forever
9 7: vboxnet2: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel
    state UP group default qlen 1000
10    inet 192.168.50.1/24 brd 192.168.50.255 scope global vboxnet2
11    valid_lft forever preferred_lft forever

```

My exposed private service is available on `http://192.168.50.240:8080/` if I want to expose this outside of the VM's I will need to add a port forwarding NAT rule on virtual box. To do this add a NAT which is on the same network as your host `10.0.0.23` in our case this will be `10.0.0.0/24`: See figure 3.

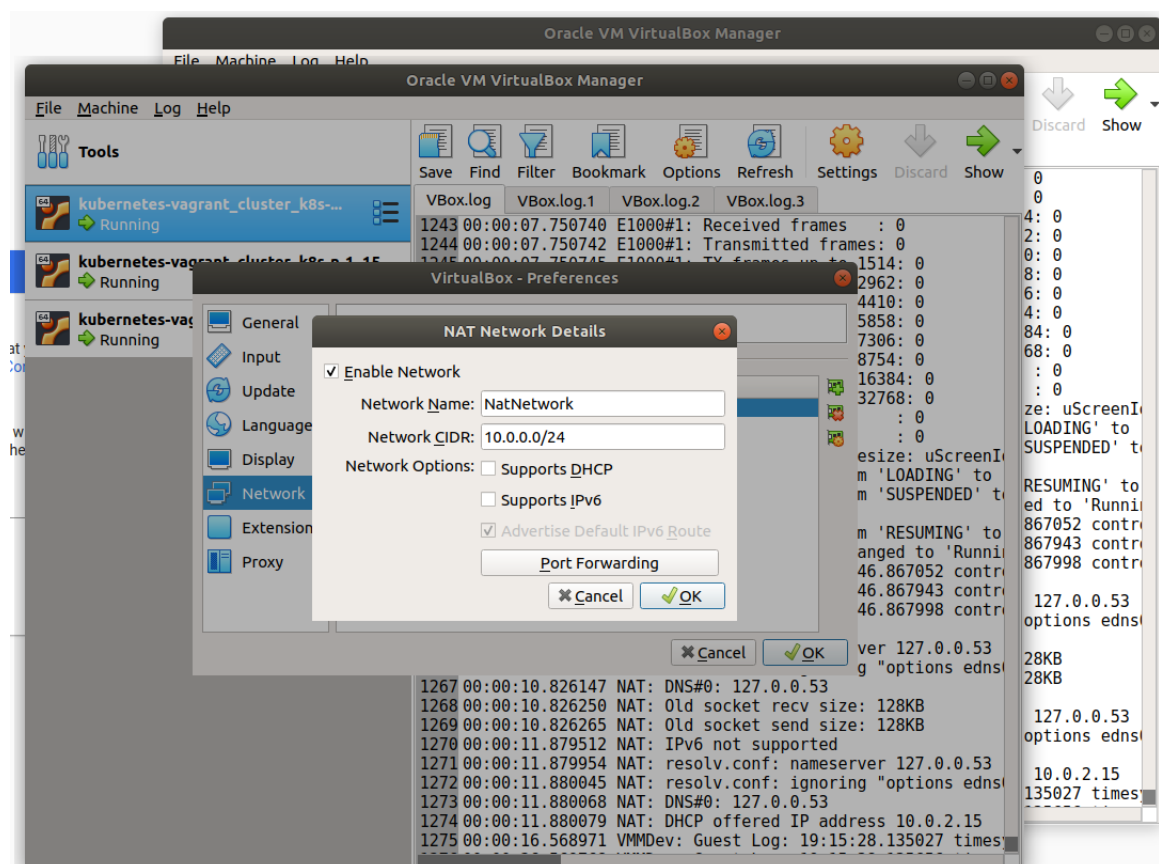


FIGURE 3: Virtualbox NAT Configuration

Add a NAT rule which maps the internal metallb external cluster IP `192.168.50.240` port `8080` and maps to the local host `10.0.0.23` port `8080`: See figure 4.

Follow this guide to add nginx proxy <https://www.hostinger.com/tutorials/how-to-set-up-nginx-reverse-proxy/>

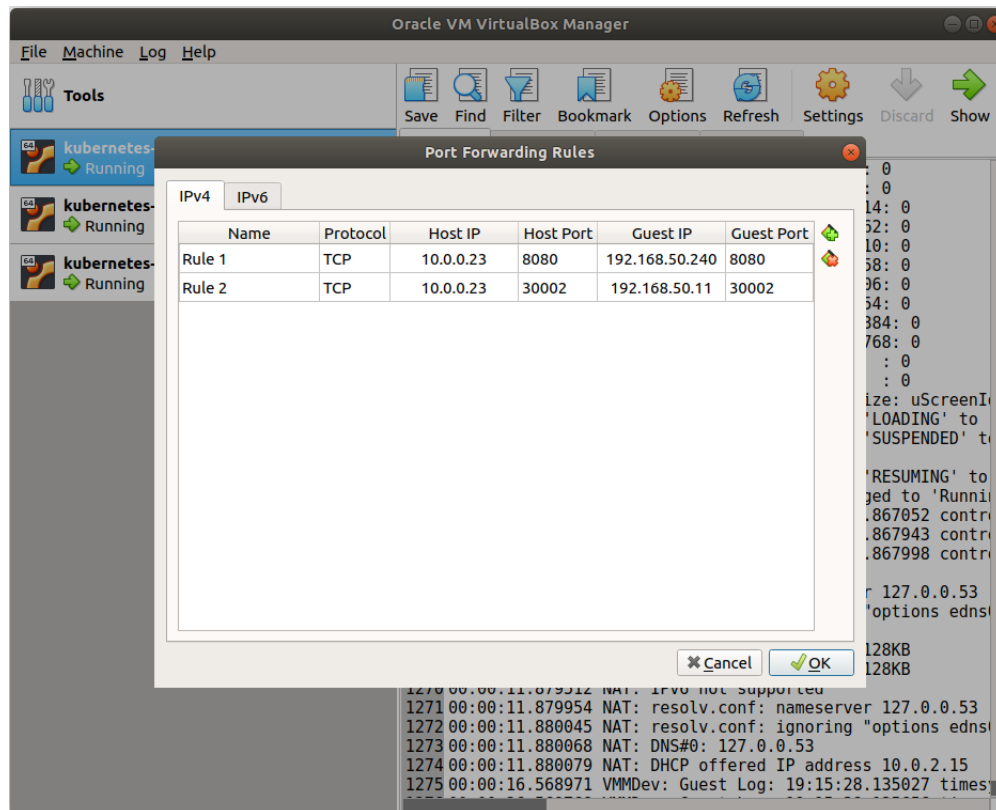


FIGURE 4: Virtualbox NAT port forwarding rule Configuration

Troubleshooting Issues:

1. Virtualbox failed with the following error:

This error means that you need to go in to the bios of your computer and enable virtualization first before running vagrant

```

1 ==> k8s-m-1: Booting VM...
2 There was an error while executing `VBoxManage`, a CLI used by Vagrant
3 for controlling VirtualBox. The command and stderr is shown below.
4
5 Command: ["startvm", "74415e44-a015-48dc-ab71-74d47a816739", "--type",
6         "headless"]
7 Stderr: VBoxManage: error: VT-x is disabled in the BIOS for all CPU
8         modes (VERR_VMX_MSR_ALL_VMX_DISABLED)
9
10 VBoxManage: error: Details: code NS_ERROR_FAILURE (0x80004005),
11        component ConsoleWrap, interface IConsole

```

Useful commands

Deploy a nginx container for quick testing of the cluster

```
1 kubectl run nginx --image nginx
```

Creating service for nginx on port 80.

- This service requires the cluster to support load balancing
- For this to work you will need a bare metal loadbalancer like [metallb](#)

```
1 kubectl expose deploy nginx --port 80 --type LoadBalancer
```

Deleting a service

```
1 kubectl delete svc nginx
```

Vagrant commands

- [General vagrant commands](#)
- [Vagrant snapshots](#)

Taking vagrant snapshots

- The following is the snapshot technique I will use when rollback is required to a good kubernetes state.
- This is especially usefull for when experiements are being run and the unexpected happens.
- Using restore on a snapshot will quickly reduce rebuild time.

```
1 vagrant snapshot save k8s-m-1 k8s-m-1--v0.3
2
3 ==> k8s-m-1: Snapshotting the machine as 'k8s-m-1--v0.1'...
4 ==> k8s-m-1: Snapshot saved! You can restore the snapshot at any time by
5 ==> k8s-m-1: using `vagrant snapshot restore`. You can delete it using
6 ==> k8s-m-1: `vagrant snapshot delete`.
7
8 vagrant snapshot save k8s-n-1 k8s-n-1--v0.3
9
10 ==> k8s-n-1: Snapshotting the machine as 'k8s-n-1--v0.1'...
```

```
11 ==> k8s-n-1: Snapshot saved! You can restore the snapshot at any time by
12 ==> k8s-n-1: using `vagrant snapshot restore`. You can delete it using
13 ==> k8s-n-1: `vagrant snapshot delete`.
14
15 vagrant snapshot save k8s-n-2 k8s-n-2--v0.3
16 ==> k8s-n-2: Snapshotting the machine as 'k8s-n-2--v0.1'...
17 ==> k8s-n-2: Snapshot saved! You can restore the snapshot at any time by
18 ==> k8s-n-2: using `vagrant snapshot restore`. You can delete it using
19 ==> k8s-n-2: `vagrant snapshot delete`.
20
21 vagrant snapshot save k8s-n-3 k8s-n-3--v0.3
22 ==> k8s-n-2: Snapshotting the machine as 'k8s-n-2--v0.1'...
23 ==> k8s-n-2: Snapshot saved! You can restore the snapshot at any time by
24 ==> k8s-n-2: using `vagrant snapshot restore`. You can delete it using
25 ==> k8s-n-2: `vagrant snapshot delete`.
```

Suspending the vagrant virtual machines

To suspend all virtual machines within the vagrant file do:

```
1 vagrant suspend
```

Expected Result

```
1 ==> k8s-m-1: Saving VM state and suspending execution...
2 ==> k8s-n-1: Saving VM state and suspending execution...
3 ==> k8s-n-2: Saving VM state and suspending execution...
```

Install Helm

Download Helm

```
1 wget https://get.helm.sh/helm-v2.16.5-linux-amd64.tar.gz
2 tar -zxvf helm-v2.16.5-linux-amd64.tar.gz
3 sudo mv linux-amd64/helm /usr/local/bin/helm
```

Create service account for tiller

```
1 kubectl -n kube-system create serviceaccount tiller
```

Create a cluster role binding for tiller

```
1 kubectl create clusterrolebinding tiller --clusterrole cluster-admin
  --serviceaccount=kube-system:tiller
```

Initilize Helm and install tiller

```
1 helm init --service-account=tiller
```

Initialize a helm chart repository

Once you have helm ready, you can add a chart repository. One popular starting location is the official Helm stable charts:

```
1 helm repo add stable https://kubernetes-charts.storage.googleapis.com/
```

you should see the following response if all is ok

```
1 "stable" has been added to your repositories
```

Install Istio []

Download Istio

```
1 curl -L https://git.io/getLatestIstio | ISTIO_VERSION=1.5.1 sh -
2 cd istio-1.5.1
3 echo 'export PATH=~/.istio-1.5.1/bin:$PATH' >> ~/.bash_profile
4 source ~/.bash_profile
5 cd ..
```

Verify istio can be installed on the cluster

```
1
2 istioctl verify-install
```

You shoow see a response like below with a passed response


```
1 Checking the cluster to make sure it is ready for Istio installation...
2
3 #1. Kubernetes-api
4 -----
5 Can initialize the Kubernetes client.
6 Can query the Kubernetes API Server.
7
8 #2. Kubernetes-version
9 -----
10 Istio is compatible with Kubernetes: v1.17.4.
11
12 #3. Istio-existence
13 -----
14 Istio will be installed in the istio-system namespace.
15
16 #4. Kubernetes-setup
17 -----
18 Can create necessary Kubernetes configurations:
19     Namespace,ClusterRole,ClusterRoleBinding,CustomResourceDefinition,Role,ServiceAccount
20
21 #5. SideCar-Injector
22 -----
23 This Kubernetes cluster supports automatic sidecar injection. To enable
24     automatic sidecar injection see
25     https://istio.io/docs/setup/kubernetes/additional-setup/sidecar-injection/#deployment
26
27 -----
28 Install Pre-Check passed! The cluster is ready for Istio installation.
```

Prepare for installing issue the following commands

Enter a Kiali username “admin” when prompted:

```
1 KIALI_USERNAME=$(read -p 'Kiali Username: ' uval && echo -n $uval |  
    base64)
```

Enter a Kiali passphrase of “admin” when prompted:

```
1 KIALI_PASSPHRASE=$(read -sp 'Kiali Passphrase: ' pval && echo -n $pval  
    | base64)
```

Make sure the istio namespace exists

```
1 NAMESPACE=istio-system  
2 kubectl create namespace $NAMESPACE
```

To create a secret, run the following commands: “‘bash cat <<EOF | kubectl apply -f - apiVersion: v1 kind: Secret metadata: name: kiali namespace: \$NAMESPACE labels: app: kiali type: Opaque data: username: \$KIALI_USERNAME passphrase: \$KIALI_PASSPHRASE EOF

Install the istio components on the cluster

```
1 istioctl manifest apply \  
2   --set profile=default \  
3   --set addonComponents.grafana.enabled=true \  
4   --set addonComponents.tracing.enabled=true \  
5   --set addonComponents.kiali.enabled=true
```

You should see the following the following response

```
1 Detected that your cluster does not support third party JWT  
    authentication. Falling back to less secure first party JWT. See  
    https://istio.io/docs/ops/best-practices/security/#configure-third-party-service-a  
    for details.  
2 - Applying manifest for component Base...  
3   Finished applying manifest for component Base.  
4 - Applying manifest for component Pilot...  
5   Finished applying manifest for component Pilot.  
6   Waiting for resources to become ready...
```

```

7   Waiting for resources to become ready...
8   Waiting for resources to become ready...
9 - Applying manifest for component IngressGateways...
10 - Applying manifest for component AddonComponents...
11   Finished applying manifest for component IngressGateways.
12   Finished applying manifest for component AddonComponents.
13
14   Installation complete

```

Determine the external IP for the gateway

```
1 kubectl get svc istio-ingressgateway -n istio-system
```

The response should look like

```

1 kubectl get svc istio-ingressgateway -n istio-system
2 NAME                                TYPE                CLUSTER-IP          EXTERNAL-IP
   PORT(S)
   AGE
3 istio-ingressgateway    LoadBalancer    10.97.104.37        192.168.50.240
   15020:32345/TCP,80:30392/TCP,443:31393/TCP,15029:31067/TCP,15030:30459/TCP,15031:30460/TCP
   13m

```

In my case I can see that the external ip address is **192.168.50.240**

Install Kubernetes Cert Manager

Installing with regular manifests All resources (the CustomResourceDefinitions, cert-manager, namespace, and the webhook component) are included in a single YAML manifest file:

Install the CustomResourceDefinitions and cert-manager itself

```

1 # Kubernetes 1.15+
2 $ kubectl apply --validate=false -f
   https://github.com/jetstack/cert-manager/releases/download/v0.14.1/cert-manager.yaml

```

Verifying the installation Once you’ve installed cert-manager, you can verify it is deployed correctly by checking the cert-manager namespace for running pods:

```
1 $ kubectl get pods --namespace cert-manager
2
3 NAME                                READY   STATUS    RESTARTS
   AGE
4 cert-manager-5c6866597-zw7kh        1/1     Running   0
   2m
5 cert-manager-cainjector-577f6d9fd7-tr771  1/1     Running   0
   2m
6 cert-manager-webhook-787858fcd-nlzsq    1/1     Running   0
   2m
```

Enable Remote access to the telemetry Addons

Follow this [guide](#) to enable external access to Graphana, Kiali, Prometheus and tracing

Visit the telemetry addons via your browser.

Kiali: `http://<IP ADDRESS OF CLUSTER INGRESS>:15029/`

Prometheus: `http://<IP ADDRESS OF CLUSTER INGRESS>:15030/`

Grafana: `http://<IP ADDRESS OF CLUSTER INGRESS>:15031/`

Tracing: `http://<IP ADDRESS OF CLUSTER INGRESS>:15032/`

Delete istio

The `uninstall` deletes the RBAC permissions, the `istio-system` namespace, and all resources hierarchically under it. It is safe to ignore errors for non-existent resources because they may have been deleted hierarchically.

```
1 $ istioctl manifest generate --set profile=demo | kubectl delete -f -
```